

BMW IDCEvo :: Ethernet driver - sxgmac/stmmac

Fixed Link for switch

The Ethernet node requires the specification on how the MAC controller is connected to the external interface. For instances where a Non-MDIO managed device is used, we need to use a fixed-link.

```
&gmac0
{
    ....
    fixed-link {
        speed = <1000>;
        full-duplex;
    };
    .....
}
```

Code Flow

Collect the data about the PHY through the platform driver

Configure the necessary in the Ethernet driver

stmmac_probe_config_dt

This sets the phylink_node as the ethernet node.

```
/* PHYLINK automatically parses the phy-handle property */
plat->phylink_node = np;
```

stmmac_dvr_probe

Needs to setup the phy connection

Here we populate the phylink struct

plcfg_link_an_mode is configured for fixed-phy.

```
stmmac_dvr_probe(*)
{
    ..
    /* Setup the phylink */
    ret = stmmac_phy_setup(priv);
    .....}

/*Configure the phylink struct */
static int stmmac_phy_setup(struct stmmac_priv *priv)
{
    struct stmmac_mdio_bus_data *mdio_bus_data = priv->plat->mdio_bus_data;
    struct fwnode_handle *fwnode = of_fwnode_handle(priv->plat->phylink_node);
    int mode = priv->plat->phy_interface; /* Refers to phy-mode(rgmii-id)
    struct phylink *phylink;

    priv->phylink_config.dev = &priv->dev->dev;
    priv->phylink_config.type = PHYLINK_NETDEV; /*Ethernet PHY */
    priv->phylink_config.pcs_poll = true;
    if (priv->plat->mdio_bus_data)
        priv->phylink_config.ovr_an_inband =
            mdio_bus_data->xpcs_an_inband;
```

```

    if (!fwnode)
        fwnode = dev_fwnode(priv->device);

    phylink = phylink_create(&priv->phylink_config, fwnode,
                            mode, &stmmac_phylink_mac_ops);
    if (IS_ERR(phylink))
        return PTR_ERR(phylink);

    if (priv->hw->xpcs)
        phylink_set_pcs(phylink, &priv->hw->xpcs->pcs);

    priv->phylink = phylink;
    return 0;
}

** Here setup a fixed phy connection
* phylink_create() - create a phylink instance
* @config: a pointer to the target &struct phylink_config
* @fwnode: a pointer to a &struct fwnode_handle describing the network
*         interface
* @iface: the desired link mode defined by &typedef phy_interface_t
* @mac_ops: a pointer to a &struct phylink_mac_ops for the MAC.
*
* Create a new phylink instance, and parse the link parameters found in @np.
* This will parse in-band modes, fixed-link or SFP configuration.
*
* Note: the rtnl lock must not be held when calling this function.
*
* Returns a pointer to a &struct phylink, or an error-pointer value. Users
* must use IS_ERR() to check for errors from this function.
*/
struct phylink *phylink_create(struct phylink_config *config,
                              struct fwnode_handle *fwnode,
                              phy_interface_t iface,
                              const struct phylink_mac_ops *mac_ops)
{
    struct phylink *pl;
    int ret;

    pl = kzalloc(sizeof(*pl), GFP_KERNEL);
    if (!pl)
        return ERR_PTR(-ENOMEM);

    mutex_init(&pl->state_mutex);
    INIT_WORK(&pl->resolve, phylink_resolve);

    pl->config = config;
    if (config->type == PHYLINK_NETDEV) {
        pl->netdev = to_net_dev(config->dev);
    } else if (config->type == PHYLINK_DEV) {
        pl->dev = config->dev;
    } else {
        kfree(pl);
        return ERR_PTR(-EINVAL);
    }

    pl->phy_state.interface = iface;
    pl->link_interface = iface;
    if (iface == PHY_INTERFACE_MODE_MOCA)
        pl->link_port = PORT_BNC;
    else
        pl->link_port = PORT_MII;
    pl->link_config.interface = iface;
    pl->link_config.pause = MLO_PAUSE_AN;
    pl->link_config.speed = SPEED_UNKNOWN;
    pl->link_config.duplex = DUPLEX_UNKNOWN;
    pl->link_config.an_enabled = true;
    pl->mac_ops = mac_ops;
    __set_bit(PHYLINK_DISABLE_STOPPED, &pl->phylink_disable_state);

```

```

timer_setup(&pl->link_poll, phylink_fixed_poll, 0);

bitmap_fill(pl->supported, __ETHTOOL_LINK_MODE_MASK_NBITS);
linkmode_copy(pl->link_config.advertising, pl->supported);
phylink_validate(pl, pl->supported, &pl->link_config);

/* Checks what type of ohy connection is present, sets pl->cfg_link_an_mode = MLO_AN_FIXED; for
fixed phy */
ret = phylink_parse_mode(pl, fwnode);
if (ret < 0) {
    kfree(pl);
    return ERR_PTR(ret);
}

/* Configure the phylink struct for fixed-phy */
if (pl->cfg_link_an_mode == MLO_AN_FIXED) {
    ret = phylink_parse_fixedlink(pl, fwnode);
    if (ret < 0) {
        kfree(pl);
        return ERR_PTR(ret);
    }
}

pl->cur_link_an_mode = pl->cfg_link_an_mode;

ret = phylink_register_sfp(pl, fwnode);
if (ret < 0) {
    kfree(pl);
    return ERR_PTR(ret);
}

return pl;
}

/* This will read the property from device tree and configure the phylink accordingly */
static int phylink_parse_fixedlink(struct phylink *pl,
                                   struct fwnode_handle *fwnode)
{
    struct fwnode_handle *fixed_node;
    const struct phy_setting *s;
    struct gpio_desc *desc;
    u32 speed;
    int ret;

    fixed_node = fwnode_get_named_child_node(fwnode, "fixed-link");
    if (fixed_node) {
        ret = fwnode_property_read_u32(fixed_node, "speed", &speed);

        pl->link_config.speed = speed;
        pl->link_config.duplex = DUPLEX_HALF;

        if (fwnode_property_read_bool(fixed_node, "full-duplex"))
            pl->link_config.duplex = DUPLEX_FULL;

        /* We treat the "pause" and "asym-pause" terminology as
         * defining the link partner's ability.
         */
        if (fwnode_property_read_bool(fixed_node, "pause"))
            __set_bit(ETHTOOL_LINK_MODE_Pause_BIT,
                    pl->link_config.lp_advertising);
        if (fwnode_property_read_bool(fixed_node, "asym-pause"))
            __set_bit(ETHTOOL_LINK_MODE_Asym_Pause_BIT,
                    pl->link_config.lp_advertising);

        if (ret == 0) {
            desc = fwnode_gpiod_get_index(fixed_node, "link", 0,
                                           GPIOD_IN, "?");

            if (!IS_ERR(desc))
                pl->link_gpio = desc;
        }
    }
}

```

```

        else if (desc == ERR_PTR(-EPROBE_DEFER))
            ret = -EPROBE_DEFER;
    }
    fwnode_handle_put(fixed_node);

    if (ret)
        return ret;
} else {
    u32 prop[5];

    ret = fwnode_property_read_u32_array(fwnode, "fixed-link",
                                         NULL, 0);

    if (ret != ARRAY_SIZE(prop)) {
        phylink_err(pl, "broken fixed-link?\n");
        return -EINVAL;
    }

    ret = fwnode_property_read_u32_array(fwnode, "fixed-link",
                                         prop, ARRAY_SIZE(prop));
    if (!ret) {
        pl->link_config.duplex = prop[1] ?
                                DUPLEX_FULL : DUPLEX_HALF;
        pl->link_config.speed = prop[2];
        if (prop[3])
            __set_bit(ETHTOOL_LINK_MODE_Pause_BIT,
                      pl->link_config.lp_advertising);
        if (prop[4])
            __set_bit(ETHTOOL_LINK_MODE_Asym_Pause_BIT,
                      pl->link_config.lp_advertising);
    }
}

if (pl->link_config.speed > SPEED_1000 &&
    pl->link_config.duplex != DUPLEX_FULL)
    phylink_warn(pl, "fixed link specifies half duplex for %dMbps link?\n",
                 pl->link_config.speed);

bitmap_fill(pl->supported, __ETHTOOL_LINK_MODE_MASK_NBITS);
linkmode_copy(pl->link_config.advertising, pl->supported);
phylink_validate(pl, pl->supported, &pl->link_config);

s = phy_lookup_setting(pl->link_config.speed, pl->link_config.duplex,
                      pl->supported, true);
linkmode_zero(pl->supported);
phylink_set(pl->supported, MII);
phylink_set(pl->supported, Pause);
phylink_set(pl->supported, Asym_Pause);
phylink_set(pl->supported, Autoneg);
if (s) {
    __set_bit(s->bit, pl->supported);
    __set_bit(s->bit, pl->link_config.lp_advertising);
} else {
    phylink_warn(pl, "fixed link %s duplex %dMbps not recognised\n",
                 pl->link_config.duplex == DUPLEX_FULL ? "full" : "half",
                 pl->link_config.speed);
}

linkmode_and(pl->link_config.advertising, pl->link_config.advertising,
             pl->supported);

pl->link_config.link = 1;
pl->link_config.an_complete = 1;

return 0;
}

```

This will connect the ethernet driver to the phylink instance

In stmmac_open we will initialize the phy connection

If an real PHY device was present "phy_attach_direct" will configure the connection and load the appropriate phydriver

In the case of fixed-phy (phylinkcfg_link_an_mode = fixed-phy), we just return true to phylink_connect. (The PHY driver is not loaded)

Start the phylink instance (phylink_start(privphylink);

```
static int stmmac_open(struct net_device *dev)
{...
if (priv->hw->pcs != STMMAC_PCS_TBI &&
    priv->hw->pcs != STMMAC_PCS_RTBI &&
    (!priv->hw->xpcs ||
     xpcs_get_an_mode(priv->hw->xpcs, mode) != DW_AN_C73)) {
    pr_alert("Harman Origin");
    /* This will return 0 as we have configured fixed-phy */
    ret = stmmac_init_phy(dev); /* This will initialize the PHY connection */
    if (ret) {
        netdev_err(priv->dev,
                    "%s: Cannot attach to PHY (error: %d)\n",
                    __func__, ret);
        goto init_phy_error;
    }
    ....
    /* Start the phy connection */
    phylink_start(priv->phylink);
}
....}

/* Check if valid PHY connection is present */
static int stmmac_init_phy(struct net_device *dev)
{
    struct stmmac_priv *priv = netdev_priv(dev);
    struct device_node *node;
    int ret;

    node = priv->plat->phylink_node;

    if (node)
    {
        pr_alert("Harman test 1");
        ret = phylink_of_phy_connect(priv->phylink, node, 0); /* this function will call
phylink_fwnode_phy_connect */
    }

    /* If the connection is a fixed-phy, then we return 0 */
int phylink_fwnode_phy_connect(struct phylink *pl,
                              struct fwnode_handle *fwnode,
                              u32 flags)
{
    struct fwnode_handle *phy_fwnode;
    struct phy_device *phy_dev;
    int ret;

    /* Fixed links and 802.3z are handled without needing a PHY */
    if (pl->q!== MLO_AN_FIXED || /* MLO_AN_FIXED checks for fixed-phy */
        (pl->cfg_link_an_mode == MLO_AN_INBAND &&
         phy_interface_mode_is_8023z(pl->link_interface)))
    {
        /* Return the phy connect status as true, this will allow for fixed-phy connections */
        return 0;
    }
}
```

Receive Process

*****Register the interrupt *****

```
stmmac_open --> ret = stmmac_request_irq(dev); --> stmmac_request_irq_multi(struct net_device *dev)
```

- Device has one common interrupt

```
ret = request_irq(dev->irq, stmmac_mac_interrupt, 0, int_name, dev);
```

- Rx AND Tx interrupts based on DMA given in eth dts(DMA 0, 12, 13, 14, 15). Rx is shown below

```
for (i = 0; i < rx_channels_count ; i++) {
22     struct stmmac_rx_queue *rx_q = &priv->rx_queue[i];
    /* basically, during probe an array is created with mapping to DMA channels. queue numbers are in
ascending order */
21     int hw_chan_idx = stmmac_map_queue_to_dma(priv, i);
20
    /* rx_irq[] is derived from dts during probe. this is the irq GIC number and +32 = IRQ number*/
19     if (priv->rx_irq[hw_chan_idx] == 0)
18         continue;
17
16     int_name = priv->int_name_rx_irq[i];
15     sprintf(int_name, "%s:%s-%d(%s-%d)", dev->name, "rxdma",
14         hw_chan_idx, "rxq", i);
13     ret = request_irq(priv->rx_irq[hw_chan_idx],
12         stmmac_rxdma_interrupt, 0, int_name, rx_q);
11     if (unlikely(ret < 0)) {
10         netdev_err(priv->dev,
9             "%s: ERROR: allocating the RxDMA IRQ %d (error: %d)\n",
8             __func__, priv->rx_irq[hw_chan_idx], ret);
7             goto rxdma_x_irq_error;
6         }
5         rxq_success_cnt++;
4     }
}
```

cat /proc/interrupts to view the created interrupts

107:	0	0	0	0	GICv3 531 Level	eth0:mac
112:	114	0	0	0	GICv3 546 Level	eth0:rxdma-0(rxq-0)
124:	0	0	0	0	GICv3 558 Level	eth0:rxdma-12(rxq-1)
125:	0	0	0	0	GICv3 559 Level	eth0:rxdma-13(rxq-2)
126:	0	0	0	0	GICv3 560 Level	eth0:rxdma-14(rxq-3)
127:	0	0	0	0	GICv3 561 Level	eth0:rxdma-15(rxq-4)
128:	0	0	0	0	GICv3 564 Level	eth0:txdma-0(txq-0)
140:	0	0	0	0	GICv3 576 Level	eth0:txdma-12(txq-1)
141:	0	0	0	0	GICv3 577 Level	eth0:txdma-13(txq-2)
142:	0	0	0	0	GICv3 578 Level	eth0:txdma-14(txq-3)
143:	2	0	0	0	GICv3 579 Level	eth0:txdma-15(txq-4)

***** Schedule the Napi*****

The Rx interrupt handler is " static irqreturn_t stmmac_rxdma_interrupt(int irq, void *data) "

The data will be the pointer to the struct stmmac_rx_queue, which has the private data.

```
5 static irqreturn_t stmmac_rxdma_interrupt(int irq, void *data)
4 {
3     struct stmmac_rx_queue *rx_q = (struct stmmac_rx_queue *)data;
2     struct stmmac_priv *priv = rx_q->priv_data;
1
0     stmmac_napi_check(priv, rx_q->queue_index, DMA_DIR_RX);
1
```

```

2         return IRQ_HANDLED;
3     }

```

This function needs to confirm the interrupt and schedule NAPI if needed

```

static int stmmac_napi_check(struct stmmac_priv *priv, u32 chan, u32 dir)

```

```

41 {
40     int status;
39     struct stmmac_rx_queue *rx_q = &priv->rx_queue[chan];
38     struct stmmac_tx_queue *tx_q = &priv->tx_queue[chan];
37     struct stmmac_channel *ch = &priv->channel[chan];
36     struct napi_struct *rx_napi;
35     struct napi_struct *tx_napi;
34     unsigned long flags;
    /* Get the dma channel from queue number */
33     int hw_chan_idx = stmmac_map_queue_to_dma(priv, chan);
32
31     rx_napi = rx_q->xsk_pool ? &ch->rxtx_napi : &ch->rx_napi;
30     tx_napi = tx_q->xsk_pool ? &ch->rxtx_napi : &ch->tx_napi;
29
28     status = stmmac_dma_interrupt_status(priv, priv->ioaddr, &priv->xstats,
27         hw_chan_idx, dir);
26
    /* Disable dma irq and Enable NAPI */
25     if ((status & handle_rx) && (chan < priv->plat->rx_queues_to_use)) {
24         if (napi_schedule_prep(rx_napi)) {
23             spin_lock_irqsave(&ch->lock, flags);
22             stmmac_disable_dma_irq(priv, priv->ioaddr, hw_chan_idx,
21                 1, 0);
20             spin_unlock_irqrestore(&ch->lock, flags);
19             __napi_schedule(rx_napi);
18         }
17     }
16
15     if ((status & handle_tx) && (chan < priv->plat->tx_queues_to_use)) {
14         if (napi_schedule_prep(tx_napi)) {
13             spin_lock_irqsave(&ch->lock, flags);
12             stmmac_disable_dma_irq(priv, priv->ioaddr, hw_chan_idx,
11                 0, 1);
10             spin_unlock_irqrestore(&ch->lock, flags);
9             __napi_schedule(tx_napi);
8         }
7     }
6
5     return status;
4 }

```